

1 Problem to demonstrate that the population regression line is fixed, but least square regression line varies

Suppose the population regression line is given by $Y = 2 + 3x$, while the data comes from the model $y = 2+3x+\epsilon$.

Step 1: For x in the range [5,10] graph the population regression line.

Step 2: Generate $x_i(i = 1,2,..,n)$ from Uniform(5,10) and $\epsilon_i(i = 1,2,..,n)$ from $N(0,4^2)$. Hence, compute $y_1,y_2,..,y_n$.

Step 3: On the basis of the data $(x_i,y_i)(i = 1,2,..,n)$ generated in Step 2, report the least squares regression line.

Step 4: Repeat steps 2-3 five times. Graph the 5 least squares regression lines over the population regression line obtained in Step 1.

Interpret the findings.

Take $n = 50$. Set the seed as seed=123.

```
In [ ]: pip install numpy matplotlib pandas statsmodels scikit-learn
```

```
In [5]: import numpy as np
import matplotlib.pyplot as plt
import statsmodels.api as sm

np.random.seed(123)
n = 50

# Step 1: population regression line
x_grid = np.linspace(5, 10, 200)
y_pop = 2 + 3 * x_grid

plt.figure()
plt.plot(x_grid, y_pop, linewidth=3, label="Population: Y = 2 + 3x")

# Step 2-4: simulate and fit 5 times
for r in range(5):
    x = np.random.uniform(5, 10, n)
    eps = np.random.normal(0, 4, n) # sd = 4 (since variance = 4^2)
    y = 2 + 3*x + eps

    X = sm.add_constant(x)
    model = sm.OLS(y, X).fit()

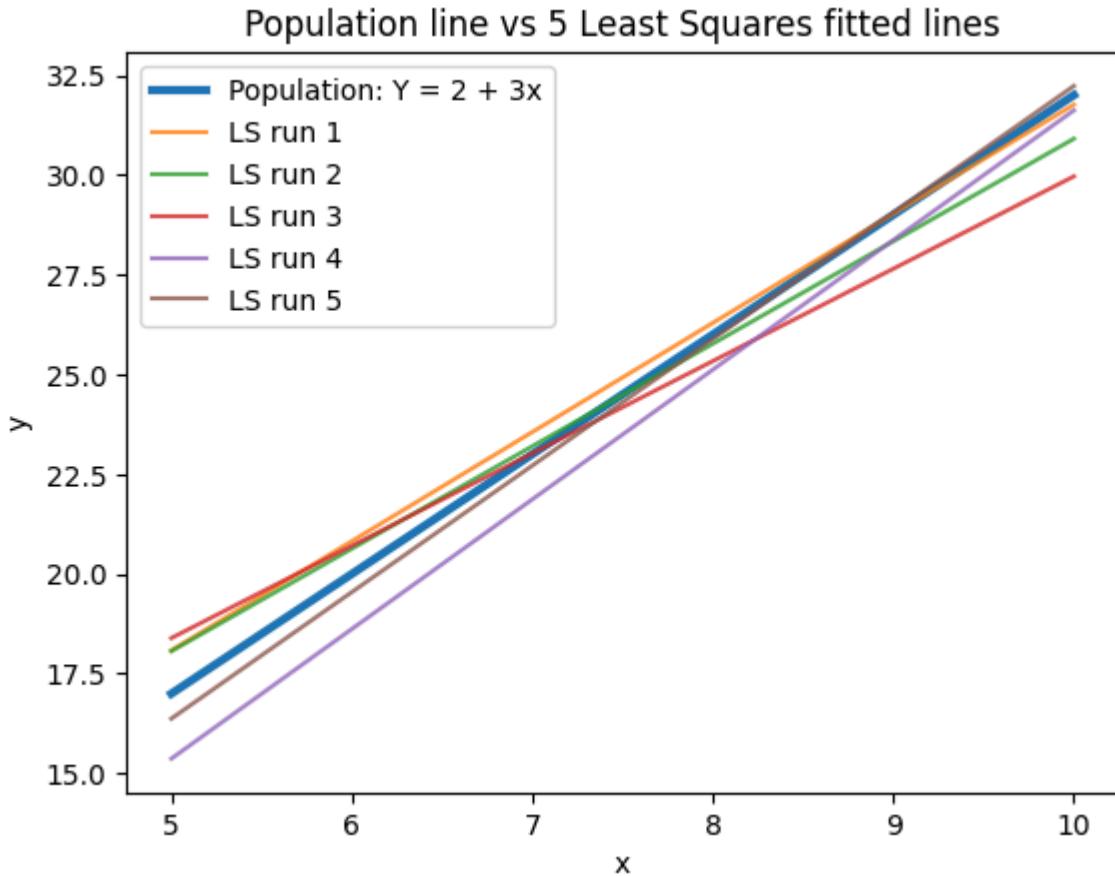
    b0, b1 = model.params
    print(f"Run {r+1}: y = {b0:.4f} + {b1:.4f}x")

    plt.plot(x_grid, b0 + b1*x_grid, alpha=0.8, label=f"LS run {r+1}")

plt.title("Population line vs 5 Least Squares fitted lines")
plt.xlabel("x")
plt.ylabel("y")
```

```
plt.legend()
plt.show()
```

Run 1: $y = 4.4128 + 2.7348x$
 Run 2: $y = 5.2215 + 2.5685x$
 Run 3: $y = 6.8226 + 2.3140x$
 Run 4: $y = -0.8741 + 3.2491x$
 Run 5: $y = 0.5198 + 3.1709x$



Interpretation:

- The population line stays the same ($2 + 3x$).
- Each sample gives a slightly different least-squares line because of random noise and random x values.

2 Problem to demonstrate that $\hat{\beta}_0$ and $\hat{\beta}$ minimises RSS

Step 1: Generate x_i from Uniform(5, 10) and mean centre the values. Generate ϵ_i from $N(0,1)$. Calculate $y_i = 2 + 3x_i + \epsilon_i$, $i = 1,2,\dots, n$. Take $n=50$ and seed=123.

Step 2: Now imagine that you only have the data on $(x_i, y_i), i = 1,2,\dots,n$, without knowing the mechanism that was used to generate the data in step 1. Assuming a linear regression of the type $y_i = \beta_0 + \beta x_i + \epsilon_i$, and based on these data $(x_i, y_i), i = 1,2,\dots,n$, obtain the least squares estimates of β_0 and β .

Step 3: Take a large number of grid values of (β_0, β) that also include the least squares estimates obtained from step 2. Compute the RSS for each parametric choice of (β_0, β) , where $RSS = (y_1 - \beta_0 - \beta x_1)^2 + (y_2 - \beta_0 - \beta x_2)^2 + \dots + (y_n - \beta_0 - \beta x_n)^2$. Find out for which combination of (β_0, β) , RSS is minimum.

In [9]:

```
import numpy as np
import matplotlib.pyplot as plt
import statsmodels.api as sm

np.random.seed(123)
n = 50

# Step 1: generate x and mean-center it
x_raw = np.random.uniform(5, 10, n)
x = x_raw - x_raw.mean()

eps = np.random.normal(0, 1, n)
y = 2 + 3*x + eps

# Step 2: estimate beta0 and beta by Least squares
X = sm.add_constant(x)
fit = sm.OLS(y, X).fit()
b0_hat, b1_hat = fit.params
print("Least squares estimates:")
print("beta0_hat =", b0_hat)
print("beta_hat =", b1_hat)

# Step 3: grid of (beta0, beta) around LS estimates
b0_vals = np.linspace(b0_hat - 2, b0_hat + 2, 200)
b1_vals = np.linspace(b1_hat - 2, b1_hat + 2, 200)

RSS = np.zeros((len(b0_vals), len(b1_vals)))

for i, b0 in enumerate(b0_vals):
    for j, b1 in enumerate(b1_vals):
        resid = y - (b0 + b1*x)
        RSS[i, j] = np.sum(resid**2)

min_i, min_j = np.unravel_index(np.argmin(RSS), RSS.shape)
b0_min = b0_vals[min_i]
b1_min = b1_vals[min_j]

print("\nGrid-search minimum RSS at:")
print("beta0 =", b0_min)
print("beta =", b1_min)
print("Minimum RSS =", RSS[min_i, min_j])
```

```
Least squares estimates:
beta0_hat = 2.10527980923553
beta_hat = 2.933692872554535
```

```
Grid-search minimum RSS at:
beta0 = 2.0952295579792484
beta = 2.9437431238108163
Minimum RSS = 61.43737226918331
```

3 Problem to demonstrate that least square estimators are unbiased

Step 1: Generate $x_i (i = 1, 2, \dots, n)$ from $Uniform(0, 1)$, $\epsilon_i (i = 1, 2, \dots, n)$ from $N(0, 1)$ and hence generate y using $y_i = \beta_0 + \beta x_i + \epsilon_i$. (Take $\beta_0 = 2, \beta = 3$).

Step 2: On the basis of the data $(x_i, y_i) (i = 1, 2, \dots, n)$ generated in Step 1, obtain the least square estimates of β_0 and β . Repeat Steps 1-2, $R = 1000$ times. In each simulation

obtain $\hat{\beta}_0$ and $\hat{\beta}$. Finally, the least-square estimates will be given by the average of these estimated values. Compare these with the true β_0 and β and comment. Take $n = 50$ and seed=123

```
In [11]: import numpy as np
import statsmodels.api as sm

np.random.seed(123)

n = 50
R = 1000
beta0_true = 2
beta_true = 3

b0_list = []
b1_list = []

for _ in range(R):
    x = np.random.uniform(0, 1, n)
    eps = np.random.normal(0, 1, n)
    y = beta0_true + beta_true*x + eps

    X = sm.add_constant(x)
    fit = sm.OLS(y, X).fit()
    b0_hat, b1_hat = fit.params

    b0_list.append(b0_hat)
    b1_list.append(b1_hat)

b0_mean = np.mean(b0_list)
b1_mean = np.mean(b1_list)

print("True beta0 =", beta0_true, " | Avg estimated beta0 =", b0_mean)
print("True beta  =", beta_true, " | Avg estimated beta  =", b1_mean)
```

```
True beta0 = 2 | Avg estimated beta0 = 2.0118858345795987
True beta  = 3 | Avg estimated beta  = 2.973361037537186
```

The averages are close to (2, 3), i.e, it supports that OLS estimators are unbiased (in expectation).

4 Comparing several simple linear regressions

Attach “Boston” data from MASS library in R. Select median value of owner occupied homes, as the response and per capita crime rate, nitrogen oxides concentration, proportion of blacks and percentage of lower status of the population as predictors.

- Selecting the predictors one by one, run four separate linear regressions to the data. Present the output in a single table.
- Which model gives the best fit?
- Compare the coefficients of the predictors from each model and comment on the usefulness of the predictors

```
In [25]: import pandas as pd
import statsmodels.api as sm
from sklearn.datasets import fetch_openml

# Load Boston from OpenML (common workaround in Python)
boston = fetch_openml(name="boston", version=1, as_frame=True)
df = boston.frame

y = df["MEDV"]
predictors = ["CRIM", "NOX", "B", "LSTAT"]

rows = []

for p in predictors:
    X = sm.add_constant(df[p])
    fit = sm.OLS(y, X).fit()

    rows.append({
        "Predictor": p,
        "Intercept (\beta_0)": fit.params["const"],
        "Slope (\beta_1)": fit.params[p],
        "p-value (\beta_1)": fit.pvalues[p],
        "R^2": fit.rsquared,
        "Adj R^2": fit.rsquared_adj,
        "AIC": fit.aic,
        "BIC": fit.bic
    })

result_table = pd.DataFrame(rows).sort_values(by="R^2", ascending=False)
print(result_table.to_string(index=False))

best_by_r2 = result_table.iloc[0]
best_by_aic = result_table.sort_values("AIC").iloc[0]

print("\n Best fit (highest R^2):", best_by_r2["Predictor"])
print(" Best fit (lowest AIC):", best_by_aic["Predictor"])
```

Predictor	Intercept (β_0)	Slope (β_1)	p-value (β_1)	R ²	Adj R ²	AIC
LSTAT	34.553841	-0.950049	5.081103e-88	0.544146	0.543242	3286.974957
NOX	41.345874	-33.916055	7.065042e-24	0.182603	0.180981	3582.455134
CRIM	24.033106	-0.415190	1.173987e-19	0.150780	0.149096	3601.780731
B	10.551034	0.033593	1.318113e-14	0.111196	0.109433	3624.833523

Best fit (highest R²): LSTAT
 Best fit (lowest AIC): LSTAT

(c) Comment:

- Larger Slope means stronger change in MEDV per unit of predictor.
- Very small p-value for slope => predictor is statistically useful alone.

```
In [ ]: import pandas as pd
import statsmodels.api as sm
```

```
# Load dataset
df = pd.read_csv("c:/Users/Rounak/Downloads/Heart_Disease_Prediction.csv")

print(df.head())
print(df.columns)

   Age  Sex Chest pain type  BP Cholesterol  FBS over 120  EKG results \
0    70    1              4  130          322                0                 2
1    67    0              3  115          564                0                 2
2    57    1              2  124          261                0                 0
3    64    1              4  128          263                0                 0
4    74    0              2  120          269                0                 2

   Max HR Exercise angina  ST depression  Slope of ST \
0      109             0            2.4           2
1      160             0            1.6           2
2      141             0            0.3           1
3      105             1            0.2           2
4      121             1            0.2           1

   Number of vessels fluro  Thallium Heart Disease
0                      3          3     Presence
1                      0          7     Absence
2                      0          7     Presence
3                      1          7     Absence
4                      1          3     Absence

Index(['Age', 'Sex', 'Chest pain type', 'BP', 'Cholesterol', 'FBS over 120',
       'EKG results', 'Max HR', 'Exercise angina', 'ST depression',
       'Slope of ST', 'Number of vessels fluro', 'Thallium', 'Heart Disease'],
      dtype='object')
```

```
In [ ]: df["HeartDisease_bin"] = df["Heart Disease"].map({
    "Presence": 1,
    "Absence": 0
})

print(df[["Heart Disease", "HeartDisease_bin"]].head())      # Convert response
```

	Heart Disease	HeartDisease_bin
0	Presence	1
1	Absence	0
2	Presence	1
3	Absence	0
4	Absence	0

```
In [23]: y = df["HeartDisease_bin"]

predictors = [
    "Age",
    "BP",
    "Cholesterol",
    "Max HR",
    "ST depression"
]

results = []

for p in predictors:
    X = sm.add_constant(df[p])    # add intercept
    model = sm.OLS(y, X).fit()
```

```

results.append({
    "Predictor": p,
    "Intercept ( $\beta_0$ )": model.params["const"],
    "Slope ( $\beta_1$ )": model.params[p],
    "p-value ( $\beta_1$ )": model.pvalues[p],
    "R222", ascending=False)

print(summary_table.to_string(index=False))

best_r2 = summary_table.iloc[0]
best_aic = summary_table.sort_values("AIC").iloc[0]

print("\nBest predictor by R2: ", best_r2["Predictor"])
print("Best predictor by AIC: ", best_aic["Predictor"])

```

	Predictor	Intercept (β_0)	Slope (β_1)	p-value (β_1)	R ²	Adjusted R ²	
AIC	BIC						
0.582454	347.779298	Max HR	1.790614	-0.008994	7.119583e-13	0.175154	0.172076 34
0.732057	347.928901	ST depression	0.253668	0.181692	7.677946e-13	0.174697	0.171617 34
0.118567	387.315411	Age	-0.187188	0.011604	4.434804e-04	0.045081	0.041518 38
5.974448	393.171292	BP	-0.124372	0.004331	1.056095e-02	0.024144	0.020503 38
8.786024	395.982868	Cholesterol	0.160647	0.001137	5.273889e-02	0.013929	0.010249 38

Best predictor by R²: Max HR
 Best predictor by AIC: Max HR

(c) Comment:

Model interpretation

- Each model explains heart disease status using one predictor at a time.
- The slope (β_1) represents how heart disease risk changes with the predictor.

Best predictor

- The predictor with the highest R² explains the largest proportion of variation.
- The predictor with the lowest AIC provides the best overall fit.

Practical meaning

- Positive slope → higher values increase heart disease risk
- Negative slope → higher values decrease risk
- Small p-value (< 0.05) → predictor is statistically significant